Contents

1 High Level View

On its core, Machine Learning about is a program that changes from exposure to data (experience) 1. Deep Learning is a subfield of machine learning, and shares this concept.


Figure 1: Machine Learning Process

Deep Learning models complex patterns of data. It's particularly useful for non-linear patterns. This opens up a new space of problems to solve. Because Deep Learning is a part of Machine Learning, the previous logic is found neural network diagrams (see Figures 2 and 3). The second image has been simplified (update isn't included).

2 The steps

There are two main steps: *forward* and *backward* propagation.

Suppose a mathematical function is given to us:

$$y(x_1, x_2) = ax_1 + bx_2 + c$$

where a, b, c are *parameters*.

In *Forward Propagation* we initialize a set of parameters (say $a, b, c = 0$, and use an equation to estimate the output (y). We input the y to another function "C" which is small if we're doing well, or large if bad.

In *Backward Propagation* we minimize C, by differentiation, and find a way to move our function parameters towards the minimum. We use *Gradient Descent*.

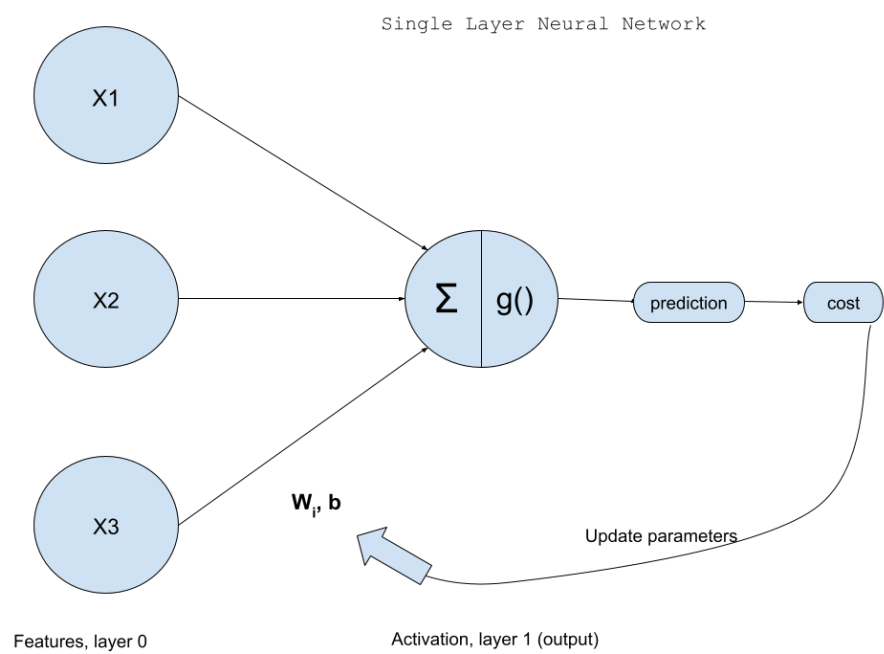Basically, we run an estimation, compute C, and use this to update our function, iteratively.

Figure 2: Single Layer Neural Network

$W^{[1](m)}_{ij}$, b_i

$W^{[2](m)}_{ij}$

Σ | g()

Σ | g()

Σ | g()

X1

X2

X3

Sum | activation

Features, layer 0

Activations, layer 1

Activation, layer 2 (output)

Figure 3: Shallow Neural Network

We will see examples in detail, starting with multivariable linear regression, that is, a previous step to *Deep Learning* and Neural Networks.

3 Example: Multivariable Linear Regression

A linear regression model for 2 features:

$$y = w_1 \cdot x_1 + w_2 \cdot x_2 + b$$
$$= \vec{w} \cdot \vec{x} + b$$

w_1, w_2 control the slopes of this plane, b translates it up and down. $\vec{x}$ is for one sample. Many samples are represented as a matrix:

$$[y1, y2] = [w1, w2] \cdot \begin{bmatrix} x_1 & x_1 \\ x_2 & x_2 \end{bmatrix} + b \tag{1}$$

y_i is the result for each sample (columns in the matrix).

We need to do forward and backward propagation.

The same than for a line (figure 4) by tuning W and b we can find the best fit. Without b the line has to go through the origin.

The sign of W reflects the line over y axis, and the magnitude controls the slope. Also b translates over x.


Figure 4: Line plot

3.1 Forward Propagation

The Loss $= Loss(w, b)$ in linear regression is the *Square Error*:

$$L_i(\vec{w}, b) = (y_i - \hat{y})^2$$
$$= (y_i - \vec{w} \cdot \vec{x}_i - b)^2$$

$\vec{w} \cdot \vec{x}_i$ can also be denoted $\sum_j w_j \cdot \mathbf{X}_{ji}$. In the first form we multiply the vector w and the column i (dot product).

In Linear Regression, the Cost is the averaged sum of L_i, and it's the *Mean Square Error*:

$$C(w_1, w_2, b) = \frac{1}{2} \sum_{i=0}^{i=2} L_i(w_1, w_2, b) \tag{2}$$

$$= \frac{1}{2}([y_1, y_2] - [\hat{y}_1, \hat{y}_2]) \cdot ([y_1, y_2] - [\hat{y}_1, \hat{y}_2]) \tag{3}$$

$$= \frac{1}{2}([y_1, y_2] - [w_1, w_2]\mathbf{X} - b) \cdot ([y_1, y_2] - [w_1, w_2]\mathbf{X} - b) \tag{4}$$

$\frac{1}{2}$ is to average over examples. Equation 4 is what we implement into code. Outliers may be critical on the effect of weights and biases.

In Python it'd look like:

```
Y' = np.dot(w,X) - b
diff = Y-Y'
cost = 1/2*np.dot(diff, diff.T)
```

3.2 Backward Propagation

We need the Cost's derivative to find better weights and bias. The complicated thing is keeping track of each step.

$$dC(\vec{w}, b) = \frac{\partial C}{\partial \vec{w}} + \frac{\partial C}{\partial b}$$

$$= d(\sum_{i=0}^{i=2} L_i) = \sum_{i=0}^{i=2} dL_i(\vec{w}, b) \tag{5}$$

$$= \sum_{i=0}^{i=2} \frac{\partial L_i}{\partial \vec{w}} + \frac{\partial L}{\partial b}$$

Equation 5 makes use of linearity property of diff. We sum over columns (each sample). Now, pick up a column/sample k:

$$L_k(\vec{w}, b) = (y_k - \hat{y_k})^2$$
$$L_k = A_k(w_1, w_2, b)^2$$
$$dL_k = 2 \cdot A_k(w1, w2, b) \cdot dA_k$$

We named A_k to the difference $y_i - \hat{y}_i$. So dC was reduced to:

$$dC = \frac{1}{2} \sum_{i=0}^{i=2} dL_i \tag{6}$$

$$= \frac{1}{2} \times 2 \cdot \sum_{i=0}^{i=2} A_i(w_1, w_2, b) \cdot dA_i \tag{7}$$

$$= \frac{1}{2} \times 2 \cdot d\mathbf{A} \cdot \mathbf{A}^T \tag{8}$$

The 2 were purposedly left, as one of those will normally be a different number, which one? This is a nice expression but we don't actually need dC. We need the partials.

We need dw and db; and for this dA_k:

$$dL_k = 2A_k \cdot \left(\frac{\partial A_k}{\partial w_1} + \frac{\partial A_k}{\partial w_2} + \frac{\partial A}{\partial b}\right)$$
$$A_k = A_k(w_1, w_2, b)$$
$$= y_k - \vec{w} \cdot \vec{x_k} - b$$
$$= y_k - w_1 \cdot x_{1k} - w_2 \cdot x_{2k} - b$$
$$dA_k = -x_{1k} - x_{2k} - 1$$

where

$$\frac{\partial A_k}{\partial w_1} = -x_{1k} \qquad \frac{\partial A_k}{\partial w_2} = -x_{2k} \qquad \frac{\partial A_k}{\partial b} = -1$$

This can be expressed in compact form:

$$dC = sum\left(-\frac{1}{2} \times 2 \cdot \mathbf{A} \cdot \mathbf{X^T} - \frac{1}{2} \times 2 \cdot \mathbf{A} \times 1\right) \tag{9}$$
$$dC = \frac{\partial C}{\partial w} + \frac{\partial C}{\partial b}$$

We don't need the sum. We also found each partial derivative.

To update the parameters, *Gradient Descent* method is used. We update the vector $\vec{w}$ as follows:

$$\vec{w} = \vec{w} - \frac{\partial C}{\partial w} \cdot \alpha \tag{10}$$

$$\vec{b} = \vec{b} - \frac{\partial C}{\partial b} \cdot \alpha \tag{11}$$

α is called *learning rate* and we use to tune the derivation. We go against the gradient so the sign is changed to $-$ (it points to max incresing direction otherwise).

4 Linear Regression: General Derivation

The problem is to find w_i, b such that the multidimensional "plane" has small error respect to each datapoint. Then for a new datapoint we will have a trained predictor.

In linear regression, the model is:

$$y = w_1 \cdot x_1 + w_2 \cdot x_2 + \ldots + w_n \cdot x_n$$

$$= \sum_i^n w_i \cdot x_i + b$$

$$= \vec{w} \cdot \vec{x} + b$$

Here $\vec{x}$ is for one sample. For n samples, it becomes a matrix, we write $\vec{y} = \vec{w} \cdot \mathbf{X} + \vec{b}$. This is represented:

$$[y_1 y_2 \ldots y_n] = [w_1 w_2 \ldots w_n] \cdot \begin{bmatrix} x_{11} & x_{12} & \ldots & x_{1m} \\ x_{21} & x_{22} & \ldots & x_{2m} \\ \vdots & \vdots & \ddots & \vdots \\ x_{n1} & x_{22} & \ldots & x_{nm} \end{bmatrix} + [b_1 b_2 \ldots b_n] \qquad (12)$$

There are m examples-columns with n features-rows. Hence $[\mathbf{X}] = m \times n$

4.1 Forward Propagation

Take the k column. The Loss $= L(w, b)$ in linear regression is the *Square Error*:

$$L_k(\vec{w}, b) = (y_k - \hat{y}_k)^2$$

$$= (y_k - \sum_{j=0}^n w_j \cdot \mathbf{X}_{jk} - b)^2$$

The cost in any method/model measures how well it's doing with he current parameters. In Linear Regression, it is the averaged sum of L_k, and it's the *Mean Square Error*:

$$C(\vec{w}, \vec{b}) = \frac{1}{m} \sum_{i=0}^m L_i(\vec{w}, b) \qquad (13)$$

$$= \frac{1}{m} \sum_{i=0}^m (\vec{y}_i - \vec{\hat{y}}_i) \cdot (\vec{y}_i - \vec{\hat{y}}_i) \qquad (14)$$

$$= \frac{1}{m} (\vec{y} - \vec{w}\mathbf{X} - \vec{b}) \cdot (\vec{y} - \vec{w}\mathbf{X} - \vec{b}) \qquad (15)$$

$C(\vec{w})$ is a way to denote C depends on all the variables in $\vec{w}$. m is the number of samples.

In Python it'd look like:

```
Y' = np.dot(w,X) - b
diff = Y-Y'
cost = 1/m*np.dot(diff, diff.T)
```

4.2 Backward Propagation

The cost is a "bowl", we will reach the global minimum (or close). We use the Cost derivative to find the better weight and bias.

$$
\begin{aligned}
dC(\vec{w}, \vec{b}) &= \frac{dC}{dw} + \frac{dC}{db} \\
&= \frac{1}{m} d(\sum_0^m L_i) \\
&= \frac{1}{m} \sum_{i=0}^m dL_i(\vec{w}, b_i) \\
&= \frac{1}{m} \sum_{i=0}^m \frac{dL_i}{d\vec{w}} + \frac{dL_i}{db} \\
&= \frac{1}{m} \sum_{i=0}^m \frac{dL_i}{dw_1} + \ldots + \frac{dL_i}{dw_n} + \frac{dL}{db}
\end{aligned}
\tag{16}
$$

Take a particular column k: Equation 16 makes use of linearity property of diff. We sum over columns (each sample). Now, pick up a column/sample k:

$$
\begin{aligned}
L_k(\vec{w}, b) &= (y_k - \hat{y}_k)^2 \\
&= A_k(w_1, w_2, \ldots, w_n, b)^2 \\
dL_k &= 2 \cdot A_k(w_1, w_2, \ldots, w_n, b) \cdot dA_k
\end{aligned}
$$

We named A_k to the difference $y_i - \hat{y}_i$.

$$
\begin{aligned}
A_k &= y_k - \sum_j w_j \cdot \mathbf{X}_{jk} - b \\
A_k &= y_k - w_1 \cdot x_{1k} \ldots - w_n \cdot x_{nk} - b
\end{aligned}
$$

$$
dA_k = \frac{\partial A}{\partial w_1} + \ldots + \frac{\partial A}{\partial w_n} + \frac{\partial A}{\partial b} dA_k = -x_{1k} - x_{2k} \ldots - x_{nk} - 1
$$

Replacing previous formulas on dC, we get:

$$dC = \frac{1}{m} \sum_{i=0}^{i=m} dL_i \tag{17}$$

$$= \frac{1}{m} \times 2 \cdot \sum_{i=0}^{i=m} A_i(w_1, w_2, b) \cdot dA_i \tag{18}$$

$$= \frac{1}{m} \times 2 \cdot d\mathbf{A} \cdot \mathbf{A}^T \tag{19}$$

This is a nice expression but we don't actually need dC. We need the partials. We need $\frac{\partial C}{\partial w}$ and $\frac{\partial C}{\partial b}$. From equation 17, it follows:

$$\frac{\partial A_k}{\partial w_1} = -x_{1k} \quad \ldots \quad \frac{\partial A_k}{\partial w_n} = -x_{nk} \qquad \frac{\partial A_k}{\partial b} = -1$$

This can be expressed in compact form:

$$dC = sum(-\frac{1}{2} \times 2 \cdot \mathbf{A} \cdot \mathbf{X^T} - \frac{1}{2} \times 2 \cdot \mathbf{A} \times 1) \tag{20}$$

$$dC = \frac{\partial C}{\partial w} + \frac{\partial C}{\partial b}$$

We don't need the sum. We also found each partial derivative.

To update the parameters, *Gradient Descent* method is used. We update the vector $\vec{w}$ as follows:

$$\vec{w} = \vec{w} - \frac{\partial C}{\partial w} \cdot \alpha \tag{21}$$

$$\vec{b} = \vec{b} - \frac{\partial C}{\partial b} \cdot \alpha \tag{22}$$

α is called *learning rate* and we use to tune the derivation. We go against the gradient so the sign is changed to $-$ (it points to max incresing direction otherwise).

There is no graphical justification as to why we update w and b like that, but we are moving w against (minus) the gradient multiplied by a constant (alpha) called *learning rate*.

The minus sign is because the gradient always points away from the minimum and we want towards it (in one dimension there are only 2 directions).

The process is called *Gradient Descent*.

Because the $y_d - y_p$ is squared, MSE is a parabola for b and w then it makes sense: the farther away we are from the minimum the larger the gradient, and the more we want to update w and b.

5   Sigmoid

The derivation is similar to Linear regression. The plot for the equation:

$$y = \frac{1}{1 + e^{-(\vec{w}\cdot\vec{x}+b)}}$$

for one dimension is on figure 5. By tuning W and b we can find the best fit.

The sign of W reflects the line over y axis, and the slope. Also b translates over x.

Figure 5: Sigmoid plot

6 Related Concepts

Overfitting: Occurs when the cost in the training dataset decreases but it increases on the test dataset. The model starts to *memorize* data, and does not *generalize*.

Training error: blue; validation error: red, both as a function of the number of training cycles. The best predictive and fitted model would be where the validation error has its global minimum.

Underfitting: It occurs when the model or algorithm does not fit the data enough. It could be a bad model (too simple, or just not the right fit), or a lack of training, etc.

Classification v Regression: A classification model is one which attempts to predict a class, or category. That is, it's predicting from a number of discrete possibilities, such as "dog" or "cat." A regression model is one which attempts to predict one or more numeric quantities, such as a temperature or a location. Which one we use depends on the nature of the variables.

Cross Validation: The goal of cross-validation is to test the model's ability to predict new data that was not used in estimating it, in order to flag problems like overfitting or selection bias, and to give an insight on how the model will generalize to an independent dataset.

Why a CNN? It's the current state-of-the-art approach to creating computer vision models.